

Query Optimisation Lab Report

1 Query Analysis

The queries will be run on the Mondial10 database. The **real** time will be used to measure query execution times, as it represents the total time from the user's perspective. It includes external delays, which can occur when queries require disk access, such delays are not captured in **user** or **sys** times.

All of the measurements are taken on an Acer Aspire A515-44, 2.38 GHz AMD Ryzen 5 with 8 GB RAM with 512 GB Flash Storage running Windows 11 V10.0.22631 and SQLite version 3.47.0.

1.1 Query Pair 1: Auto Index Optimisation

SQLite is able to optimise queries using automatic indices, which can significantly improve performance by avoiding full table scans during lookups or joins. In this query pair, one will be optimisable where SQLite can create an index, while the other will not be optimisable due to the use of transformations on indexed columns.

The queries will select cities with a population greater than 1,000,000, along with the country that the city belongs to.

1.1.1 Query 1A: Optimisable by using automatic index

```
SELECT c.Name, ci.Name
FROM Country c
JOIN City ci ON ci.Country = c.Code
WHERE ci.Population > 1000000;
```

This query would be optimisable by SQLite, as the optimiser would create an auto index on the country table when joining the city table.

1.1.2 Query 1B: Transformation on JOIN Column - Not Optimisable

```
SELECT c.Name, ci.Name
FROM Country c
JOIN City ci ON CAST(ci.Country AS TEXT) = CAST(c.Code AS TEXT)
WHERE ci.Population > 1000000;
```

This query is similar to 1A, except that both columns in the JOIN are type-cast to **TEXT**.

1.1.3 Hypothesis

SQLite should optimise query 1A by creating an automatic index on the `Country.Code` column for the JOIN and possibly on `City.Population` for filtering. This should result in an efficient query plan.

In query 1B, because of the use of `CAST` in the JOIN condition, SQLite should not be able to create any indices on `Country.Code` or `City.Country`. Indices only work on raw column values, and so this query should not be optimisable since the `CAST` function transforms the two columns. Instead, it should perform full table scans on both tables, leading to a less efficient query plan with higher execution times.

1.1.4 Results and Discussion

Query	Run 1 (s)	Run 2 (s)	Run 3 (s)	Run 4 (s)	Run 5 (s)	Average (s)
Query 1A	0.067	0.067	0.065	0.066	0.065	0.0660
Query 1B	0.806	0.805	0.800	0.814	0.808	0.8066

Table 1: Execution Times on Mondial10 Database

Query Execution Plan of 1A:

```
|--SCAN ci
'--SEARCH c USING AUTOMATIC COVERING INDEX (Code=?)
```

Query Execution Plan of 1B:

```
|--SCAN ci
'--SCAN c
```

Table 1 shows that query 1A’s average execution time (0.0660) was around 91.8% faster than query 1B (0.8066). Query 1A benefited from SQLite’s use of automatic covering indices, with an index on `Country.Code` to allow searching.

In contrast, query 1B relied on full table scans due to type-cast transformations on the JOIN columns, preventing index usage. Indices are more efficient as they reduce the search space, allowing SQLite to directly locate relevant rows and minimise comparisons and disk I/O. Full table scans (as shown in 1B’s execution plan), however, process all rows sequentially, which is costly in a large database. The `CAST` operation forces SQLite to treat the columns as derived expressions, making them ineligible for indexing.

1.2 Query Pair 2: Subquery Optimisation

This query pair explores how SQLite optimises nested subqueries, comparing a derived table approach and a correlated subquery. Both queries aim to retrieve the names of countries alongside the average population of their cities, but their performance depends significantly on how SQLite’s optimiser processes them.

1.2.1 Query 2A: Optimisable by Query Flattening or Co-Routines

```
SELECT c.Name, s.AvgPop AS AverageCityPopulation
FROM Country c, (
    SELECT ci.Country, AVG(ci.Population) AS AvgPop
    FROM City ci
    GROUP BY ci.Country
) s
WHERE c.Code = s.Country;
```

This query calculates the average population of cities by grouping the `City` table in a subquery within the `FROM` clause. The result is treated as a derived table (`s`) and joined with the `Country` table.

1.2.2 Query 2B: Correlated Subquery - Not Optimisable

```
SELECT c.Name,
    (SELECT AVG(Population)
     FROM City
     WHERE City.Country = c.Code) AS AverageCityPopulation
FROM Country c;
```

Query 2B uses a correlated subquery within the `SELECT` clause. For every row in the `Country` table, the subquery calculates the average city population by scanning the `City` table.

1.2.3 Hypothesis

The SQLite optimiser can optimise query 2A through query flattening or co-routine execution. Query flattening transforms the subquery into a `JOIN`, eliminating the need for separate execution and enabling more efficient data retrieval. Alternatively, SQLite may use co-routines, which process the subquery in parallel with the outer query, delivering rows incrementally as needed. Both techniques should enhance performance.

In contrast, query 2B’s correlated subquery structure should hinder optimisation. The subquery must be evaluated separately for each row of the `Country` table, as it references values

from the outer query. This dependency should lead to multiple scans of the `City` table, increasing execution time in comparison to query 2A.

1.2.4 Results and Discussion

Query	Run 1 (s)	Run 2 (s)	Run 3 (s)	Run 4 (s)	Run 5 (s)	Average (s)
Query 2A	0.059	0.061	0.059	0.061	0.059	0.0598
Query 2B	5.266	5.245	5.248	5.306	5.353	5.2836

Table 2: Execution Times on Mondial10 Database

Query Execution Plan of 2A:

```
--CO-ROUTINE s
|  |--SCAN ci
|  '--USE TEMP B-TREE FOR GROUP BY
|--SCAN c
'--SEARCH s USING AUTOMATIC COVERING INDEX (Country=?)
```

Query Execution Plan of 2B:

```
--SCAN c
'--CORRELATED SCALAR SUBQUERY 1
  '--SCAN City
```

From the results in Table 2, we can observe that query 2A performed significantly better than 2B. There was approximately a 98% decrease in execution time from query 2B to 2A.

The query plans show that the SQLite optimiser was able to effectively utilise co-routines in query 2A to run the subquery once in parallel with the outer query. This allows the outer query to process rows as they are needed without waiting for the entire subquery to complete. The optimiser also uses a temporary B-Tree to optimise aggregation (`GROUP BY`), and uses an automatic covering index to optimise joins between the subquery and the outer query.

In the execution plan of 2B, the optimiser is not able to use parallelism or indexing. The correlated subquery in the `SELECT` clause makes the subquery dependent on the current row of the `Country` table, forcing it to scan the entire `City` table for each row in `Country`. This explains the increased execution time observed.

1.3 Query Pair 3: JOIN Order Optimisation

This query pair examines SQLite’s join reordering optimisation, using different join strategies to retrieve countries with populations over 10,000,000, their capitals, and continents they belong to, with continent area being over 30,000,000 km².

1.3.1 Query 3A: JOIN - Optimisable by Reordering

```
SELECT Country.Name, City.Name, City.Population, Continent.Name
FROM Country
JOIN encompasses ON encompasses.Country = Country.Code
JOIN Continent ON encompasses.Continent = Continent.Name
JOIN City ON City.Name = Country.Capital AND City.country = Country.Code
WHERE Continent.Area > 30000000 AND Country.Population > 10000000;
```

This query involves (INNER) JOINS between the `Country`, `encompasses`, `Continent`, and `City` tables.

1.3.2 Query 3B: CROSS JOIN - Order Not Optimisable

```
SELECT Country.Name, City.Name, City.Population, Continent.Name
FROM City
CROSS JOIN encompasses
CROSS JOIN Continent
CROSS JOIN Country
WHERE Continent.Area > 30000000 AND Country.Population > 10000000
      AND encompasses.Country = Country.Code
      AND encompasses.Continent = Continent.Name
      AND City.Name = Country.Capital
      AND City.country = Country.Code;
```

In this query, CROSS JOINS are used between the `City`, `encompasses`, `Continent`, and `Country` tables. The filtering conditions in the `WHERE` clause are applied to ensure that valid rows are returned, and matches results from query 3A.

1.3.3 Hypothesis

SQLite should optimise the query 3A by selecting an efficient join order through join reordering. Since (INNER) JOINS are commutative, the optimiser can reorder the joins to find the most efficient execution order. The query planner should evaluate the cost of potential join orders, considering statistics and factors such as the table sizes, available indices, and filtering conditions, and choose an optimal join order.

On the other hand, query 3B uses CROSS JOINS, which do not allow join reordering. This means that the tables must be processed in the order they appear, the optimiser would not be able to adjust the order to minimise work. This approach should lead to larger intermediate results and require more processing time, resulting in a less efficient query plan compared to query 3A.

1.3.4 Results and Discussion

Query	Run 1 (s)	Run 2 (s)	Run 3 (s)	Run 4 (s)	Run 5 (s)	Average (s)
Query 3A	1.180	1.162	1.195	1.178	1.184	1.1798
Query 3B	7.543	7.161	7.088	7.102	7.263	7.2314

Table 3: Execution Times on Mondial10 Database

Query Execution Plan of 3A:

```
|--SCAN Country
|--SEARCH encompasses USING AUTOMATIC COVERING INDEX (Country=?)
|--SEARCH Continent USING AUTOMATIC PARTIAL COVERING INDEX (Name=?)
'--SEARCH City USING AUTOMATIC COVERING INDEX (Name=? AND Country=?)
```

Query Execution Plan of 3B:

```
|--SCAN City
|--SCAN encompasses
|--SEARCH Continent USING AUTOMATIC PARTIAL COVERING INDEX (Name=?)
'--SEARCH Country USING AUTOMATIC PARTIAL COVERING INDEX (Code=? AND Capital=?)
```

Query 3A’s average execution time (1.1798) is around 83% faster than 3B, which takes approximately 6.1 seconds longer. Query 3A benefits from join reordering, which allows SQLite to explore efficient join sequences based on table sizes, filtering conditions, and available indexes. It is also able to scan one table and create indices for the subsequent joins. The indices allow searching, which will only process relevant rows, instead of processing all rows in a scan. This results in significantly faster execution, as the intermediate data processed at each step is minimised.

Query 3B, on the other hand, suffers from the rigidity of **CROSS JOINS**, which don’t allow reordering. The query generates a Cartesian product of all rows from the tables, leading to larger intermediate results. It does however use automatic indices for the filtering in the **Country** and **Continent** table, this is probably influenced by optimising the filters in the **WHERE** clause, which are similar to those that would appear in the **ON** clause of a **JOIN**.

2 Reflection

Figure 1 illustrates the impact of SQLite’s query optimisation techniques across the three scenarios. All of the optimised queries show a clear reduction in time in comparison to their unoptimised counterparts. The index usage and subquery optimisation queries show near-instant execution with times less than 0.1 seconds. This is because they directly reduce




Figure 1: Query Pairs Execution Time Bar Plot

redundant operations, such as scanning the same table multiple times or evaluating intermediate results unnecessarily. The difference in the join reordering queries is relatively modest compared to the other scenarios. The number of joins involved could explain the greater execution times.

The experimentation also emphasises the need for developers to write queries that align with database optimisation strategies. Avoiding pitfalls like correlated subqueries ensures scalability, particularly with a large database, as seen in query 2B. Similarly, leveraging `JOINS` and understanding join reordering allows query planners to minimise intermediate results, improving performance, as demonstrated in query 3A. Developers should also design queries to utilise indexing effectively, avoiding transformations on indexed columns, as highlighted in query 1B. By adopting these best practices, developers can reduce execution times, optimise resource usage, and ensure their applications perform efficiently, even with complex queries and extensive data.

However, the experiment has limitations in fully accounting for optimisation performance, highlighting the need for tailored query design. All of the queries were run on the Mondial10 database, smaller or larger databases could amplify or reveal shortcomings of these optimisations. For example, there might not be a huge difference between using and not using an automatic index on a smaller database. Furthermore, the optimiser uses heuristics and

statistics about the database to choose an optimisation strategy, this may lead to varying optimisation strategies depending on database sizes. Uneven or skewed data distributions further complicate performance, sparse tables may weaken index efficiency, while clustered or repetitive data might amplify indexing benefits but reduce join reordering effectiveness.